

Portable Systems Group

NT OS/2 IRP Language Definition

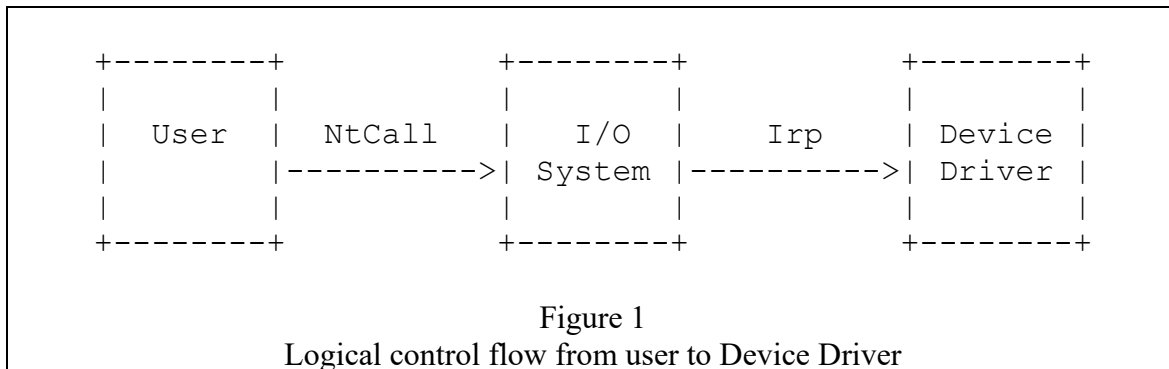
Author: *Gary D. Kimura*

Revision 1.0x, December 15, 1989

1. Introduction.....	1
2. Valid IRP combination	2
2.1 Disk Driver IRPs.....	2
2.2 File System IRPs.....	3
2.3 Keyboard Driver IRPs.....	5
2.4 Mouse Driver IRPs	5
2.5 Network Drivers IRPs.....	6
2.6 Sound Driver IRPs	6
2.7 Tape Driver IRPs	6
2.8 Terminal Driver IRPs.....	6
2.9 Video Driver IRPs.....	6
3. IRP Function Descriptions.....	8
3.1 Close	8
3.2 Create	11
3.3 Device Control	19
3.4 Directory Control(Notify Change Directory)	19
3.5 Directory Control(Query Directory)	19
3.6 File System Control(Dismount Volume)	19
3.7 File System Control(Lock Volume).....	19
3.8 File System Control(Mount Volume)	19
3.9 File System Control(Query Information File System).....	19
3.10 File System Control(Set Information File System)	19
3.11 File System Control(Unlock Volume)	19
3.12 File System Control(Verify Volume)	19
3.13 Internal Device Control.....	19
3.14 Lock Control(Lock)	19
3.15 Lock Control(Unlock All).....	19
3.16 Lock Control(Unlock Single)	20
3.17 Query Acl.....	20
3.18 Query Ea	20
3.19 Query Information	20
3.20 Query Volume Information.....	20
3.21 Read	20
3.22 Read Terminal.....	20
3.23 Set Acl.....	20
3.24 Set Ea	20
3.25 Set Information	20
3.26 Set New Size	20
3.27 Set Volume Information	20
3.28 Write	20

1. Introduction

The purpose of this chapter is to define the semantic contents of an I/O Request Packet (IRP). The information contained here is intended for use mainly by Device Driver and File System developers. The I/O system sends to the various Device Drivers¹ a stream of multiple IRPs that the drivers must interpret and respond to. Figure 1 shows the relationship between the device driver and the I/O system. Communication between the I/O system and the Device Driver is through IRPs. This chapter concentrates on the IRP language.



Each IRP has a well defined format and semantic meaning, and the order in which they are sent must adhere to certain rules. The ordering of IRPs and responses form a context sensitive language.

Each IRP contains a common header section followed by one or more function specific records (also called IRP stack locations). From a Device Drivers viewpoint each IRP request is a single record describing one function to perform. That is, the drivers only interpret one function specific record. The additional stack locations are for use when a driver issues subsequent IRPs to a lower level driver and wishes to reuse the original IRP.

Each IRP function is identified by a major and minor function field in the IRP stack location record. The list of possible function combinations are listed below. Each line lists a major function code followed (in paranthesis) by a minor function code. Note that some major functions (e.g., CREATE) do not make use the minor function field.

```

CLOSE()
CONFIGURATION_CONTROL(...)
CREATE()
DEVICE_CONTROL(...)
DIRECTORY_CONTROL(NOTIFY_CHANGE_DIRECTORY)
DIRECTORY_CONTROL(QUERY_DIRECTORY)
FILE_SYSTEM_CONTROL(DISMOUNT_VOLUME)
  
```

¹For clarity we will use the term Device Driver to refer to both Device Drivers and File systems.

```

FILE_SYSTEM_CONTROL(LOCK_VOLUME)
FILE_SYSTEM_CONTROL(MOUNT_VOLUME)
FILE_SYSTEM_CONTROL(QUERY_INFO_FILE_SYSTEM)
FILE_SYSTEM_CONTROL(SET_INFO_FILE_SYSTEM)
FILE_SYSTEM_CONTROL(UNLOCK_VOLUME)
FILE_SYSTEM_CONTROL(VERIFY_VOLUME)
INTERNAL_DEVICE_CONTROL(...)
LOCK_CONTROL(LOCK)
LOCK_CONTROL(UNLOCK_ALL)
LOCK_CONTROL(UNLOCK_SINGLE)
QUERY_ACL()
QUERY_EA()
QUERY_INFORMATION()
QUERY_VOLUME_INFORMATION()
READ()
READ_TERMINAL()
SET_ACL()
SET_EA()
SET_INFORMATION()
SET_NEW_SIZE()
SET_VOLUME_INFORMATION()
WRITE()

```

```

/* We need to define the minor function codes for the configuration, device, and internal
device function codes. */

```

Each Device Driver will only receive a combination of the preceding function codes based on the drivers device type. This means that a file system device driver can expect to receive different functions than the keyboard device driver, or a disk driver. The possible device driver types are:

```

Disk Driver,
File System (including network redirector),
Keyboard Driver,
Mouse Driver,
Network Drivers,
Sound Driver,
Tape Driver,
Terminal Driver, and
Video Driver,

```

```

/* We will need to futher expand on the different network device drivers */

```

The remainder of this chapter describes the valid combination of IRP function codes that each different device driver can expect to receive. This is followed by a section listing every IRP function code along with a description of the function's parameters, semantics, and I/O completion status codes.

2. Valid IRP Function Combinations

The section contains an individual table for each device driver type that lists the set of valid IRP functions that can be sent to the driver and under what conditions the functions are sent.

2.1 Disk Driver IRPs

The set of possible IRPs that can be sent to a disk driver are:

<u>IRP Function</u>	<u>When sent</u>
CLOSE	Anytime.
CREATE	Anytime.
DEVICE_CONTROL (...)	Anytime.
READ	Anytime.
WRITE	Anytime.

2.2 File System IRPs

The set of possible IRPs that can be sent to a file system are:

<u>IRP Function</u>	<u>When sent</u>
CLOSE	Only after a successful CREATE and then only on an opened file. This closes the file so no other operation can be performed on the file other than CREATE.
CREATE	Only after a successful MOUNT_VOLUME and then only on a mounted volume that is not locked. If successful the file is considered opened.
DIRECTORY_CONTROL (NOTIFY_CHANGE_DIRECTORY)	Only after a successful CREATE and then only on an opened directory file.
DIRECTORY_CONTROL (QUERY_DIRECTORY)	Only after a successful CREATE and then only on an opened directory file.
FILE_SYSTEM_CONTROL (DISMOUNT_VOLUME)	Only after a successful MOUNT_VOLUME and then only on a mounted volume. This

	dismounts the volume, so no other operation can be performed on the volume other than MOUNT_VOLUME.
FILE_SYSTEM_CONTROL (LOCK_VOLUME)	Only after a successful CREATE and then only on an opened file. This locks the volume containing the file such that no other creates using the same volume will succeed until the volume is unlocked. To be successful, the file used to lock the volume must also be the only opened file on the volume.
FILE_SYSTEM_CONTROL (MOUNT_VOLUME)	Anytime. If the operation is successful then a new device object for the volume is created and the volume is considered mounted and not locked.
FILE_SYSTEM_CONTROL (QUERY_INFO_FILE_SYSTEM)	Only after a successful CREATE and then only on an opened file.
FILE_SYSTEM_CONTROL (SET_INFO_FILE_SYSTEM)	Only after a successful CREATE and then only on an opened file.
FILE_SYSTEM_CONTROL (UNLOCK_VOLUME)	Only after a successful CREATE and then only on a opened file. The file system must handle the situation where the user is attempting to unlock a volume that is not locked. If successful this operation unlocks a previously locked volume so that other creates using the volume can now succeed.
FILE_SYSTEM_CONTROL (VERIFY_VOLUME)	Only after a successful MOUNT_VOLUME and then only on a mounted volume.
LOCK_CONTROL (LOCK)	Only after a successful CREATE and then only on an opened file. If successful this operation locks a range of bytes within a file. The locks remain in affect until they are explicitly unlocked or the file is closed.
LOCK_CONTROL (UNLOCK_ALL)	Only after a successful CREATE and then only on an opened file. The file system must handle the situation where an unlock is received even though there are no outstanding locks for that user.

LOCK_CONTROL (UNLOCK_SINGLE)	Only after a successful CREATE and then only on an opened file. The file system must handle the situation where an unlock is received even though there is not a corresponding lock.
QUERY_ACL	Only after a successful CREATE and then only on an opened file.
QUERY_EA	Only after a successful CREATE and then only on an opened file.
QUERY_INFORMATION	Only after a successful CREATE and then only on an opened file.
QUERY_VOLUME_INFORMATION	Only after a successful CREATE and then only on an opened file.
READ	Only after a successful CREATE and then only on an opened file.
SET_ACL	Only after a successful CREATE and then only on an opened file.
SET_EA	Only after a successful CREATE and then only on an opened file.
SET_INFORMATION	Only after a successful CREATE and then only on an opened file.
SET_NEW_SIZE	Only after a successful CREATE and then only on an opened file.
SET_VOLUME_INFORMATION	Only after a successful CREATE and then only on an opened file.
WRITE	Only after a successful CREATE and then only on an opened file.

2.3 Keyboard Driver IRPs

The set of possible IRPs that can be sent to the Keyboard driver are:

<u>IRP Function</u>	<u>When sent</u>
---------------------	------------------

CLOSE	Anytime.
CREATE	Anytime.
DEVICE_CONTROL (...)	Anytime.
QUERY_INFORMATION	Anytime.
READ	Anytime.
SET_INFORMATION	Anytime.
WRITE	Anytime.

2.4 Mouse Driver IRPs

The set of possible IRPs that can be sent to the Mouse driver are:

<u>IRP Function</u>	<u>When sent</u>
CLOSE	Anytime.
CREATE	Anytime.
DEVICE_CONTROL (...)	Anytime.
QUERY_INFORMATION	Anytime.
READ	Anytime.
SET_INFORMATION	Anytime.
WRITE	Anytime.

2.5 Network Drivers IRPs

The set of possible IRPs that can be sent to the Network drivers are:

<u>IRP Function</u>	<u>When sent</u>
---------------------	------------------

/* This table needs to be filled in */

2.6 Sound Driver IRPs

The set of possible IRPs that can be sent to the Sound driver are:

IRP Function	When sent
/* This table needs to be filled in */	

2.7 Tape Driver IRPs

The set of possible IRPs that can be sent to the Tape driver are:

IRP Function	When sent
/* This table needs to be filled in */	

2.8 Terminal Driver IRPs

The set of possible IRPs that can be sent to the Terminal driver are:

IRP Function	When sent
/* This table needs to be filled in */	

2.9 Video Driver IRPs

The set of possible IRPs that can be sent to the Video driver are:

IRP Function	When sent
CLOSE	Anytime.
CREATE	Anytime.
DEVICE_CONTROL (...)	Anytime.
QUERY_INFORMATION	Anytime.
READ	Anytime.
SET_INFORMATION	Anytime.
WRITE	Anytime.

3. IRP Function Descriptions

This section describes the input parameters and semantics for each IRP function code. It also discusses the interactions between the parameters and lists possible return status codes.

The parameter descriptions list all the fields that are used within the IRP by the operation being described. Each parameter is either Read (i.e., used as input to the operation), Set (i.e., used as output for the operation), or Ignored. To help distinguish the parameters we will also use the two terms *IrpFlags* and *FunctionFlags* to denote the flags field of the IRP header and the I/O stack location respectively.

In the description of the return status codes we do not include generic values such as `STATUS_PENDING` or `STATUS_INVALID_PARAMETER` which can be returned for any IRP. We also do not describe values that can be returned by a lower level device drivers such as `STATUS_PARITY_ERROR`.

3.1 Close

The close function is used to close a previously opened file or directory. Its two input parameters are a device object and an IRP. The device object parameter points to a volume previously mounted by the Device Driver and is where the file opened file exists. The IRP contains the close function parameters (and are listed below).

Besides closing the file, this function will optionally deletes the file based upon the disposition specified by the caller (See the `SET_INFORMATION` operation). If this is the last file object with the file opened and the disposition is *delete on close* then the file is removed from the on-disk structure.

```
Close (
    IN PDEVICE_OBJECT DeviceObject,
    IN PIRP Irp
);
```

Parameters within the IRP:

<u>Parameter type and name</u>	<u>Description</u>
PMDL <i>MdlAddress</i>	Ignored.
ULONG <i>IrpFlags</i>	Ignored.
STRING <i>FileObject->FileName</i>	Ignored.

ULONG*FileObject->RelatedFileObject*

Ignored.

PVOID*FileObject->FsContext*

Read and Set. The driver uses this field to retrieve any private data (established by the CREATE function) that needs to be processed in order to close the file. It is set to NULL upon return from the close function.

PVOID*FileObject->FsContext2*

Read and Set. The driver uses this field to retrieve any private data (established by the CREATE function) that needs to be processed in order to close the file. It is set to NULL upon return from the close function.

PVOID*FileObject->SectionObjectPointer*

Set. The close function must set this field to NULL.

IO_STATUS_BLOCK*IoStatus*

Set. This receives the final return status of the operation. The possible return status values are listed later.

PEPROCESS*AlternateProcess*

Ignored.

KPROCESSOR_MODE*RequestorMode*

Ignored.

PVOID*SystemBuffer*

Ignored.

PIO_STATUS_BLOCK*UserIoSb*

Ignored.

PKEVENT*UserEvent*

Ignored.

LARGE_INTEGER*AllocationSize*

Ignored.

PVOID*UserBuffer*

Ignored.

Parameters within the IRP Stack:

<u>Parameter type and name</u>	<u>Description</u>
UCHAR <i>MajorFunction</i>	Read. Must be equal to IRP_MJ_CLOSE.
UCHAR <i>MinorFunction</i>	Ignored.
UCHAR <i>FunctionFlags</i>	Ignored.
UCHAR <i>Control</i>	Ignored.

Iosb Return Status and Information:

The following status codes are used to complete the CLOSE function.

<u>Return status followed by information field of IOSB</u>	<u>Description</u>
STATUS_SUCCESS Ignored	Indicates that the opened file has been closed.

3.2 Create

The create function is used to create or open a file or a directory. Its two input parameters are a device object and an IRP. The device object parameter points to a volume previously mounted by the Device Driver and is where the file will exist. The IRP contains the create function parameters (and are listed below).

```
Create (
    IN PDEVICE_OBJECT DeviceObject,
    IN PIRP Irp
);
```

Parameters within the IRP:

<u>Parameter type and name</u>	<u>Description</u>
PMDL <i>MdlAddress</i>	Ignored.
ULONG <i>IrpFlags</i>	Ignored.
STRING <i>FileObject->FileName</i>	Read. This is the name of the file being opened.
ULONG <i>FileObject->RelatedFileObject</i>	Read. This field is used for path relative file names. If it is null then the file name is relative to the root of the volume (e.g., "\CONFIG.SYS" is the name of the configuration file located in root directory). If it is not null then it points to a previously opened file object representing a directory on the volume, and the file name is relative to the specified directory (e.g., if the related file object is "\NT\SDK" the file name can be "INC\NTIOAPI.H"). Note that path relative file names do not begin with a backslash.
PVOID <i>FileObject->FsContext</i>	Set. This is used by the Device Driver to store file object specific information that can be retrieved later when the driver is called to perform subsequent operations on the file.

PVOID*FileObject->FsContext2*

The FAT file system stores in this field a pointer to an internal File Control Block (FCB) structure.

Set. This is used by the Device Driver to store file object specific information that can be retrieved later when the driver is called to perform subsequent operations on the file.

The FAT file system only uses this field for directories. It is a pointer to an internal Context Control Block (CCB) structure.

PVOID*FileObject->SectionObjectPointer*

Set. It is set to the longword context for the file. It is not used for directories. For every opened file the driver allocates a single longword of context for exclusive use by the memory management system. All file objects that denote the same file point to the same longword context.

In FAT this is done by reserving a longword field in the FCB and having each section object pointer point to this field.

IO_STATUS_BLOCK*IoStatus*

Set. This receives the final return status of the operation. The possible return status values are listed later.

PEPROCESS*AlternateProcess*

Ignored.

KPROCESSOR_MODE*RequestorMode*

Read. This is the mode of the requestor. It is used for to help decide if the requestor has the proper access rights to the file.

/**** We also need to pass in the token of the requestor ****/

PVOID*SystemBuffer*

Read. This field is only used if the file is being created and then it only specifies the optional extended attributes for the file. If the field is null the file will not be created with extended attributes. The create operation must complete with an error if there are any

problems with the extended attributes.

For FAT there is a 64K limit to the size of the extended attributes (as packed on the disk). The create operation will complete with an error if this limit is exceeded.

PIO_STATUS_BLOCK

UserIoSb

Ignored.

PKEVENT

UserEvent

Ignored.

LARGE_INTEGER

AllocationSize

Read. This field is only used if the file is being created and is ignored for directories and for open operations. It specifies the initial file allocation in bytes to allocate to the file. This is not the same as the end-of-file location.

PVOID

UserBuffer

Ignored.

Parameters within the IRP Stack:

<u>Parameter type and name</u>	<u>Description</u>
UCHAR <i>MajorFunction</i>	Read. Must be equal to IRP_MJ_CREATE.
UCHAR <i>MinorFunction</i>	Ignored.
UCHAR <i>FunctionFlags</i>	Ignored.
UCHAR <i>Control</i>	Ignored.
ULONG <i>DesiredAccess</i>	Read. This is the access mask that the user is trying to acquire to the file. If the user is trying to open a file the mask will be a combination of the following values:

DELETE,
 READ_CONTROL,
 WRITE_DAC,
 WRITE_OWNER,
 SYNCHRONIZE,
 FILE_READ_DATA,
 FILE_WRITE_DATA,
 FILE_APPEND_DATA,
 FILE_READ_EA,
 FILE_WRITE_EA,
 FILE_EXECUTE,
 FILE_READ_ATTRIBUTES, and
 FILE_WRITE_ATTRIBUTES.

If the user is trying to open a directory the mask will be a combination of the following values:

DELETE,
 READ_CONTROL,
 WRITE_DAC,
 WRITE_OWNER,
 SYNCHRONIZE,
 FILE_LIST_DIRECTORY,
 FILE_ADD_FILE,
 FILE_ADD_SUBDIRECTORY,
 FILE_READ_EA,
 FILE_WRITE_EA,
 FILE_TRAVERSE,
 FILE_DELETE_CHILD,
 FILE_READ_ATTRIBUTES, and
 FILE_WRITE_ATTRIBUTES.

The driver must ensure that the combination of the caller's privileges and requestor's mode grants all of the desired accesses that the user is trying to acquire.

ULONG

Options

Read. This field contains all of the different create options and create disposition flags that the user can specify in an NT call. The valid flags and their meanings are listed below:

FILE_CREATE_DIRECTORY

Read. Indicates that the user is creating a

	new directory.
FILE_OPEN_DIRECTORY	Read. Indicates that the user is opening an existing directory.
FILE_WRITE_THROUGH	Ignored, but saved away for use by subsequent read and write operations to the file object.
FILE_SEQUENTIAL_ONLY	Ignored, but saved away for use by subsequent read and write operations to the file object.
FILE_MAPPED_IO	Ignored, but saved away for use by subsequent read and write operations to the file object.
FILE_DISABLE_CACHING	Ignored, but saved away for use by subsequent read and write operations to the file object.
FILE_SYNCHRONOUS_IO_ALERT	Ignored.
FILE_SYNCHRONOUS_IO_NONALERT	Ignored.
FILE_CREATE_TREE_CONNECTION	Read. Only used by the network. /**** need a complete description of this parameter ****/
FILE_SUPERSEDE << 24 ²	Read. Indicates that if the file already exists it should be superseded, and if the file does not exist it should be created.
FILE_CREATE << 24	Read. Indicates that if the file already exists it is an error, and if the file does not exist it should be created.
FILE_OPEN << 24	Read. Indicates that if the file already exists it is to be opened, and if the file

²To test if the flags FILE_SUPERSEDE, FILE_OPEN, FILE_CREATE, and FILE_OPEN_IF are in the options parameter the driver must first shift the flag 24 bits to the left and then do the test (e.g., Option & (FILE_SUPERSEDE << 24)).

FILE_OPEN_IF << 24

does not exist it is an error.

Read. Indicates that if the file already exists it is to be opened, and if the file does not exist it should be created.

/**** We need a list of the illegal flag combinations, and state that they will never be seen in an IRP ****/

USHORT

FileAttributes

Read. This field specifies the DOS file attributes to use when creating or superseding a file, and is ignored when opening an existing file. It is a combination of any of the following flags:

FILE_ATTRIBUTE_READONLY
FILE_ATTRIBUTE_HIDDEN
FILE_ATTRIBUTE_SYSTEM
FILE_ATTRIBUTE_ARCHIVE
FILE_ATTRIBUTE_CONTROL, and
FILE_ATTRIBUTE_NORMAL

The flag FILE_ATTRIBUTE_NORMAL overrides all other file attribute flags. (i.e., if the user specifies normal and readonly then the file is created as a normal file and not readonly).

USHORT

ShareAccess

Read. This field specifies the share mode access between processes trying to open the same file. All users that open a file for shared access must specify the exact same share flags. This is separate from their desired access. For example a file opened shared read, write, and delete, must be opened by all users as shared read, write, and delete even though the desired access might only specify read access.

The valid flags and their meanings are listed below:

FILE_SHARE_READ

Read. Indicates that the file can be opened by others for read access. If the

	file is already opened for shared read access then other users can open it for read access.
FILE_SHARE_WRITE	Read. Indicates that the file can be opened by others for write access. If the file is already opened for shared write access then other users can open it for write access.
FILE_SHARE_DELETE	Read. Indicates that the file can be opened by others for delete access. If the file is already opened for shared delete access then other users can open it for delete access.
FILE_SHARE_RENAME	Read. Indicates that the file can be renamed by others. If the file is already opened for shared renamed access then other users can rename the file.
	The test that a user requesting shared read, write, or delete can be done by the Device Driver during the create operation (i.e., a user is allowed read access to a shared file if the shared access flags match, shared read is specified, and the file's security protection allows for read access). The test for rename access must be deferred until the a rename IRP is processed (see the Set Information IRP description).
ULONG <i>EaLength</i>	Read. This parameter is specified only if the user is creating or superseding a file and has specified an EA for the file. This parameter is then the size, in bytes, of the EA set specified by the user. (i.e., it is the size of the system buffer parameter).

Iosb Return Status and Information:

The following status codes are used to complete the CREATE function.

Return status followed by

<u>information field of IOSB</u>	<u>Description</u>
STATUS_SUCCESS FILE_OPENED	Indicates that an existing file has been successfully located and opened.
STATUS_SUCCESS FILE_SUPERSEDED	Indicates that an existing file has been successfully located and superseded.
STATUS_SUCCESS FILE_CREATED	Indicates that an existing file (of the same name) does not exist and that a new file has been successfully created.
STATUS_ACCESS_DENIED Ignored	Indicates that because of protection on the file, parent directory, or volume access has been denied to the file. This can also occur if the caller specified options or share access flags are not compatible with either the file or the previous share access that it was opened with.
STATUS_OBJECT_NAME_INVALID Ignored	Indicates that the last name in the object's file name field does not contain a syntactically valid name (e.g., it's too long or contains invalid characters).
STATUS_OBJECT_NAME_NOT_FOUND Ignored	Indicates that the last name in the object's file name field is not the name of an existing file.
STATUS_OBJECT_PATH_INVALID Ignored	Indicates that a name within the path part of the object's file name field does not contain a syntactically valid name.
STATUS_OBJECT_PATH_NOT_FOUND Ignored	Indicates that a name within the path part of the object's file name field does not contain the name of an existing directory.
STATUS_DISK_FULL_ERROR Ignored	Indicates that because the disk is full the file cannot be created. This can occur when disk space cannot be allocated for the directory entry, file node, or the extended attributes.
STATUS_DISK_FULL_WARNING FILE_SUPERSEDED	Indicates that the file has been superseded but because the disk is full the file cannot be given the user specified file allocation size.
STATUS_DISK_FULL_WARNING	Indicates that the file has been created but because

FILE_CREATED

the disk is full the file cannot be given the user specified file allocation size.

STATUS_EA_INVALID
Ignored

Indicates that the EA structure passed into this function is syntactically invalid.

3.3 Device Control

3.4 Directory Control(Notify Change Directory)

3.5 Directory Control(Query Directory)

3.6 File System Control(Dismount Volume)

3.7 File System Control(Lock Volume)

3.8 File System Control(Mount Volume)

The mount function is used mount a new disk volume. Its two input parameters are a device object and an IRP. The device object parameter points to the Device Drivers original device object that is created when the driver is initialized.

The mount operation can handle mounting new volume, and remounting a previously mounted volume. The parameter description that follows assumes that it is processing a new volume. At the end of the description we cover the updating required for the remount case.

```
Mount (
    IN PDEVICE_OBJECT DeviceObject,
    IN PIRP Irp
);
```

Parameters within the IRP:

<u>Parameter type and name</u>	<u>Description</u>
PMDL <i>MdlAddress</i>	Ignored.
ULONG <i>IrpFlags</i>	Ignored.
PFILE_OBJECT <i>FileObject</i>	Ignored.
IO_STATUS_BLOCK <i>IoStatus</i>	Set. This receives the final return status of the operation. The possible return status values are listed later.
PEPROCESS <i>AlternateProcess</i>	Ignored.
KPROCESSOR_MODE <i>RequestorMode</i>	Ignored.
PVOID <i>SystemBuffer</i>	Ignored.
PIO_STATUS_BLOCK <i>UserIoCb</i>	Ignored.
PKEVENT	Ignored.

*UserEvent***LARGE_INTEGER***AllocationSize*

Ignored.

PVOID*UserBuffer*

Ignored.

Parameters within the IRP Stack:

<u>Parameter type and name</u>	<u>Description</u>
UCHAR <i>MajorFunction</i>	Read. Must be equal to IRP_MJ_FILE_SYSTEM_CONTROL.
UCHAR <i>MinorFunction</i>	Read. Must be equal to IRP_MN_MOUNT_VOLUME.
UCHAR <i>FunctionFlags</i>	Ignored.
UCHAR <i>Control</i>	Ignored.
PDEVICE_OBJECT <i>Vpb->DeviceObject</i>	Set. If the mount is successful this field is set the point to the newly allocated device object for the volume. If the mount is unsuccessful or this is a remount then this field is not updated.
ULONG <i>Vpb->DeviceObject->Flags</i>	Set. If the mount is successful then the flag DO_DIRECT_IO is set in the newly created device objects flags field. Setting this flag allows the Device Driver to receive unbuffered I/O requests for this volume.
ULONG <i>Vpb->SerialNumber</i>	Set. If the mount is successful this field is set to the serial number found on the volume. It is ignored if the mount is unsuccessful or in the case of a remount.
CHAR <i>Vpb->VolumeName[20]</i>	Set. If the mount is successful this field is set to the label found on the volume. If the volume does not have a label then this field

should be set to all spaces.

For FAT the volume label, if present, is found in the root directory as a special dirent.

PDEVICE_OBJECT
DeviceObject

Read. This is the device object that the Device Driver is to use when formulating IRPs to read or write to the volume. It is also called the target device object. If the volume is mounted successful this value must be remembered so the driver can handle subsequent requests to the volume.

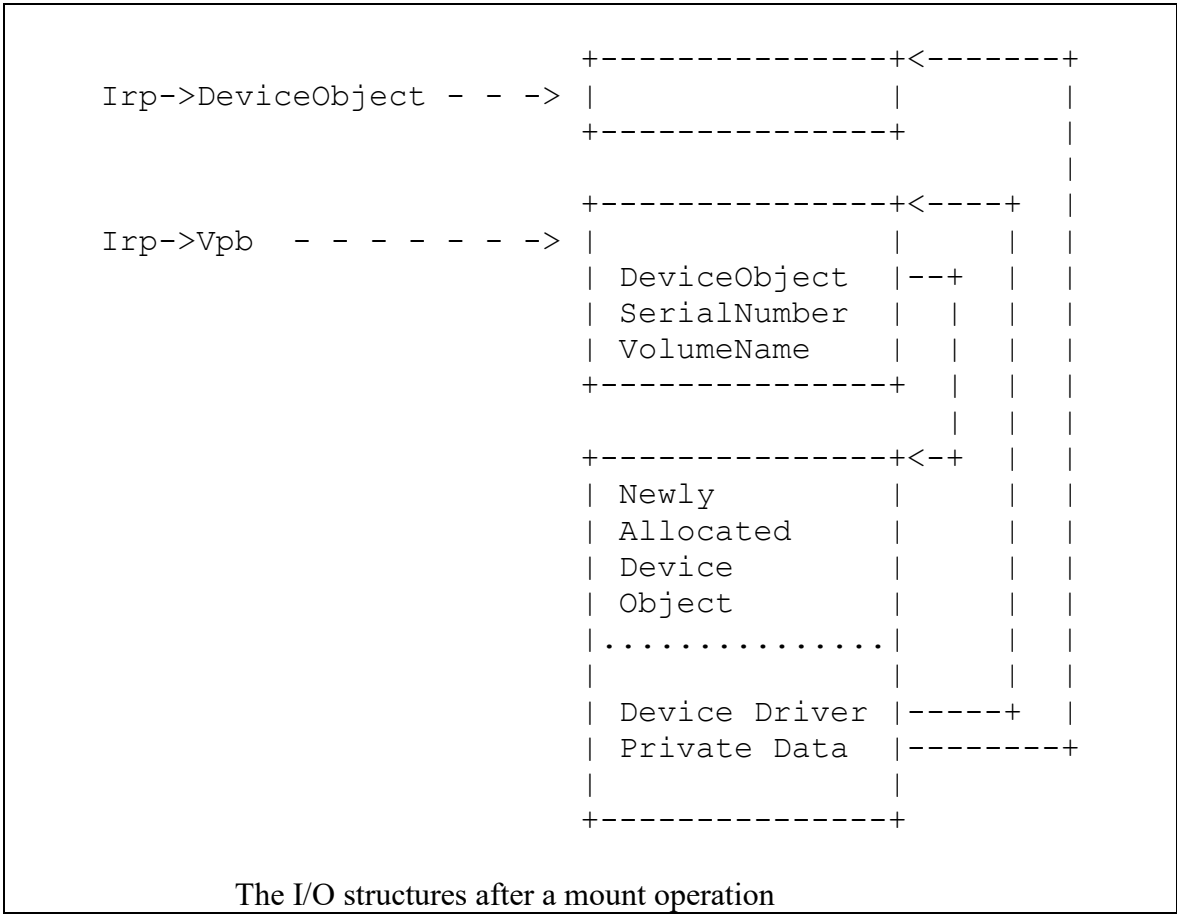
Iosb Return Status and Information:

The following status codes are used to complete the MOUNT function.

<u>Return status followed by information field of IOSB</u>	<u>Description</u>
STATUS_SUCCESS Ignored	Indicates that the volume has been successful mounted.
STATUS_WRONG_VOLUME Ignored	Indicates that the volume cannot be mounted either because it does not recognize the on-disk structure or the on-disk structure has been corrupted.

Mounting a new volume:

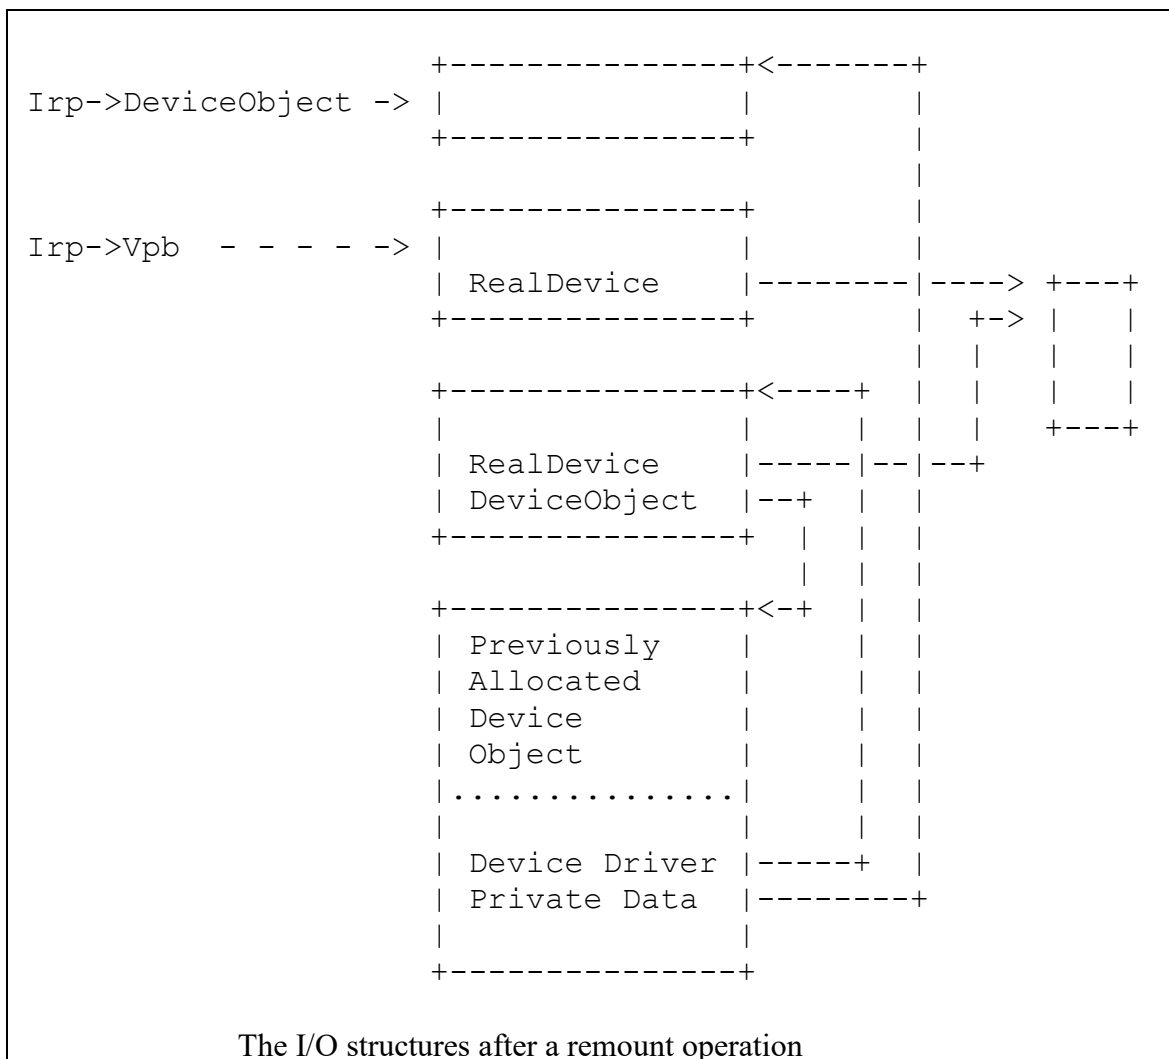
The following figure shows the major I/O structures after processing a successful mount request.



In the preceding figure the newly allocated device object has immediately following it a Device Driver private data record that is for used only by the driver. This technique should be used in the driver to keep track of the VPB and the device object where it is to send its read and write requests. It should also be used to link together all of the mounted volumes serviced by the driver.

Remounting a volume:

By using the device driver private data record to maintain a link of all mounted volumes a Device Driver can determine if a mount request for a volume matches a previously mounted volume (They match if the both volume have the same serial number and volume label). The following figure shows the major I/O structure after processing a remount.



The remount operation does not allocate any new structures, instead it performs the following operations:

- o The Device Drivers Private Data pointer to the target device object is changed to point to the new target device object.
- o The RealDevice field of the Vpb that we previously mounted is set to the RealDevice field of the new Vpb that was passed in as a parameter in the IRP.
- o The Irp->Vpb is deallocated from pool by the device driver, and complete the mount request with STATUS_SUCCESS.

3.9 File System Control(Query Information File System)

3.10 File System Control(Set Information File System)

3.11 File System Control(Unlock Volume)

3.12 File System Control(Verify Volume)

3.13 Internal Device Control

3.14 Lock Control(Lock)

3.15 Lock Control(Unlock All)

3.16 Lock Control(Unlock Single)

3.17 Query Acl

3.18 Query Ea

3.19 Query Information

3.20 Query Volume Information

3.21 Read

3.22 Read Terminal

3.23 Set Acl

3.24 Set Ea

3.25 Set Information

3.26 Set New Size

3.27 Set Volume Information

3.28 Write

Revision History

Original Draft 1.0, December 15, 1989